



Bus-based computer systems




Agenda

- Input and output
- Busses
- Memory Architectures



I/O devices

- Usually includes some non-digital component.
- Typical digital interface to CPU:




Application: 8251 UART

- Universal asynchronous receiver transmitter (UART) : provides serial communication.
- 8251 functions are integrated into standard PC interface chip.
- Allows many communication parameters to be programmed.



Serial communication

- Characters are transmitted separately:




Serial communication parameters

- Baud (bit) rate.
- Number of bits per character.
- Parity/no parity.
 - Even/odd parity.
- Length of stop bit (1, 1.5, 2 bits).



825 I CPU interface



Programming I/O

- Two types of instructions can support I/O:
 - special-purpose I/O instructions;
 - memory-mapped load/store instructions.
- Intel x86 provides in, out instructions. Most other CPUs use memory-mapped I/O.
- I/O instructions do not preclude memory-mapped I/O.



ARM memory-mapped I/O

- Define location for device:

```
DEVI EQU 0x1000
```

- Read/write code:

```
LDR r1,#DEVI ; set up device adrs
```

```
LDR r0,[r1] ; read DEVI
```

```
LDR r0,#8 ; set up value to write
```

```
STR r0,[r1] ; write value to device
```



Peek and poke

- Traditional HLL interfaces:

```
int peek(char *location) {  
    return *location; }
```

```
void poke(char *location, char newval) {  
    (*location) = newval; }
```



Busy/wait output

- Simplest way to program device.
 - Use instructions to test when device is ready.

```
current_char = mystring;  
while (*current_char != '\0') {  
    poke(OUT_CHAR,*current_char);  
    while (peek(OUT_STATUS) != 0);  
    current_char++;  
}
```



Simultaneous busy/wait input and output

```
while (TRUE) {  
    /* read */  
  
    while (peek(IN_STATUS) == 0);  
    achar = (char)peek(IN_DATA);  
  
    /* write */  
  
    poke(OUT_DATA, achar);  
    poke(OUT_STATUS, 1);  
    while (peek(OUT_STATUS) != 0);  
}
```

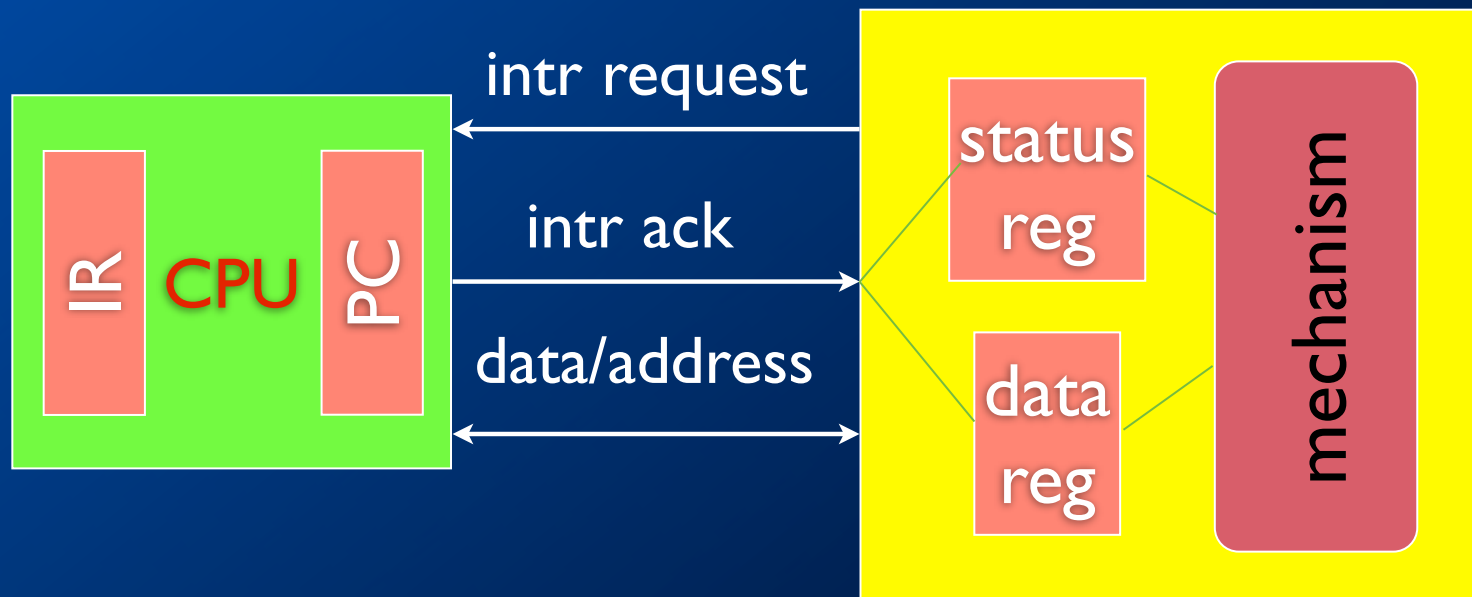


Interrupt I/O

- Busy/wait is very inefficient.
 - CPU can't do other work while testing device.
 - Hard to do simultaneous I/O.
- Interrupts allow a device to change the flow of control in the CPU.
 - Causes subroutine call to handle device.



Interrupt interface



Interrupt behavior

- Based on subroutine call mechanism.
- Interrupt forces next instruction to be a subroutine call to a predetermined location.
- Return address is saved to resume executing foreground program.



Interrupt physical interface

- CPU and device are connected by CPU bus.
- CPU and device handshake:
 - device asserts interrupt request;
 - CPU asserts interrupt acknowledge when it can handle the interrupt.



Example: character I/O handlers

```
void input_handler() {  
    achar = peek(IN_DATA);  
    gotchar = TRUE;  
    poke(IN_STATUS,0);  
}  
  
void output_handler() {  
}
```




Example: interrupt-driven main program

```
main() {  
    while (TRUE) {  
        if (gotchar) {  
            poke(OUT_DATA,achar);  
            poke(OUT_STATUS,I);  
            gotchar = FALSE;  
        }  
    }  
    other processing....  
}
```



Example: interrupt I/O with buffers

- Queue for characters:



leave one empty slot to allow full buffer to be detected

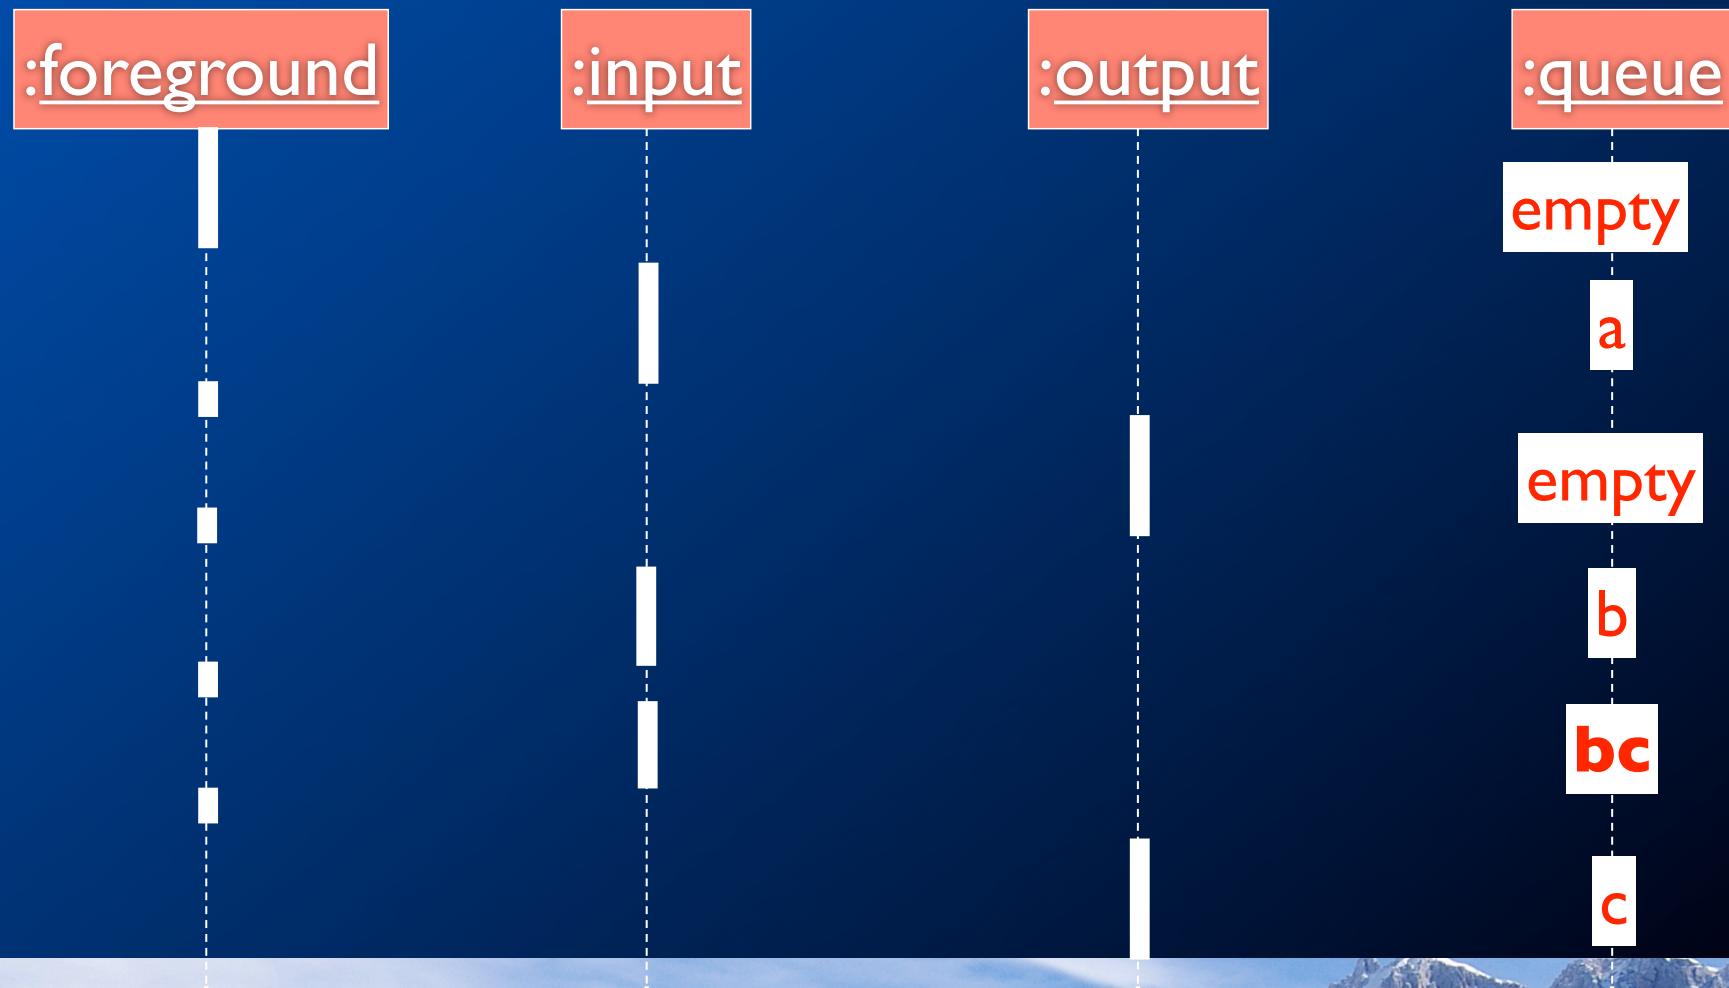


Buffer-based input handler

```
void input_handler() {  
    char achar;  
    if (full_buffer()) error = 1;  
    else {  
        achar = peek(IN_DATA);  
        add_char(achar);  
        poke(IN_STATUS,0);  
        if (nchars == 1)  
        {  
            poke(OUT_DATA,remove_char());  
            poke(OUT_STATUS,1);}  
    }  
}
```



I/O sequence diagram



Debugging interrupt code

- What if you forget to change registers?
 - Foreground program can exhibit mysterious bugs.
 - Bugs will be hard to repeat---depend on interrupt timing.




Priorities and vectors

- Two mechanisms allow us to make interrupts more specific:
 - **Priorities** determine what interrupt gets CPU first.
 - **Vectors** determine what code is called for each type of interrupt.
- Mechanisms are orthogonal: most CPUs provide both.



Prioritized interrupts




Interrupt prioritization

- **Masking**: interrupt with priority lower than current priority is not recognized until pending interrupt is complete.
- **Non-maskable interrupt (NMI)**: highest-priority, never masked.
 - Often used for power-down.




Example: Prioritized I/O




Interrupt vectors

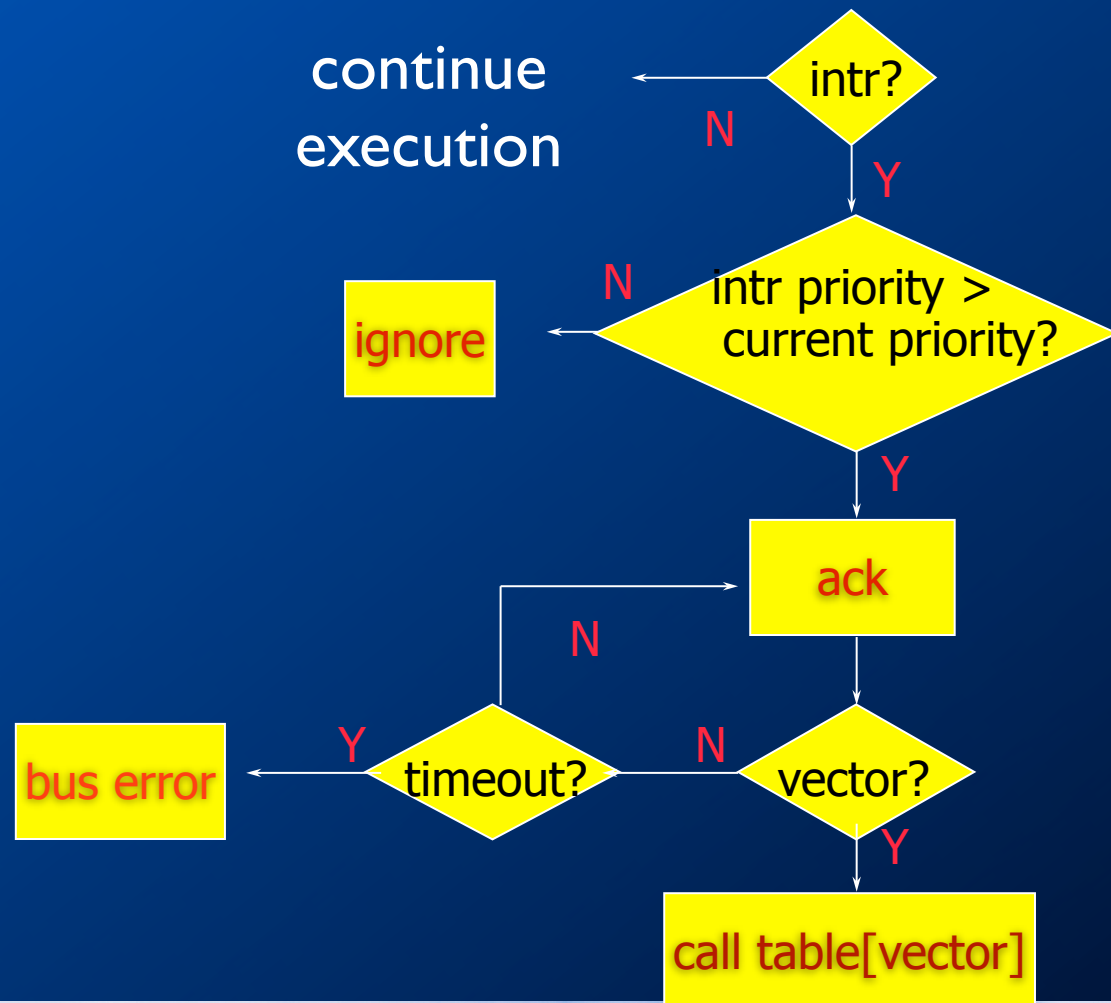
- Allow different devices to be handled by different code.
- Interrupt vector table:



Interrupt vector acquisition



Generic interrupt mechanism



Assume priority selection is handled before this point.



Interrupt sequence

- CPU acknowledges request.
- Device sends vector.
- CPU calls handler.
- Software processes request.
- CPU restores state to foreground program.



Sources of interrupt overhead

- Handler execution time.
- Interrupt mechanism overhead.
- Register save/restore.
- Pipeline-related penalties.
- Cache-related penalties.



ARM interrupts

- ARM7 supports two types of interrupts:
 - Fast interrupt requests (FIQs).
 - Interrupt requests (IRQs).
- Interrupt table starts at location 0.



ARM interrupt procedure

- CPU actions:
 - Save PC. Copy CPSR to SPSR.
 - Force bits in CPSR to record interrupt.
 - Force PC to vector.
- Handler responsibilities:
 - Restore proper PC.
 - Restore CPSR from SPSR.
 - Clear interrupt disable flags.



ARM interrupt latency

- Worst-case latency to respond to interrupt is 27 cycles:
 - Two cycles to synchronize external request.
 - Up to 20 cycles to complete current instruction.
 - Three cycles for data abort.
 - Two cycles to enter interrupt handling state.



Supervisor mode

- May want to provide protective barriers between programs.
- Avoid memory corruption.
- Need **supervisor mode** to manage the various programs.



ARM supervisor mode

- Use SWI instruction to enter supervisor mode, similar to subroutine:
 - SWI CODE_I
- Sets PC to 0x08.
- Argument to SWI is passed to supervisor mode code.
- Saves CPSR in SPSR.



Exception

- Exception: internally detected error.
- Exceptions are synchronous with instructions but unpredictable.
- Build exception mechanism on top of interrupt mechanism.
- Exceptions are usually prioritized and vectorized.



Trap

- **Trap (software interrupt)**: an exception generated by an instruction.
- Call supervisor mode.
- ARM uses SWI instruction for traps.



Co-processor

- **Co-processor**: added function unit that is called by instruction.
- Floating-point units are often structured as co-processors.
- ARM allows up to 16 designer-selected co-processors.
 - Floating-point co-processor uses units 1, 2.



Agenda

- Input and output
- Busses
- Memory Architectures



The CPU bus

- Bus allows CPU, memory, devices to communicate.
 - Shared communication medium.
- A bus is:
 - A set of wires.
 - A communications protocol.




Bus protocols

- Bus protocol determines how devices communicate.
- Devices on the bus go through sequences of states.
 - Protocols are specified by state machines, one state machine per actor in the protocol.
- May contain asynchronous logic behavior.




Four-cycle handshake



device 1

device 2




Four-cycle handshake, cont'd.

1. Device 1 raises enq.
2. Device 2 responds with ack.
3. Device 2 lowers ack once it has finished.
4. Device 1 lowers enq.

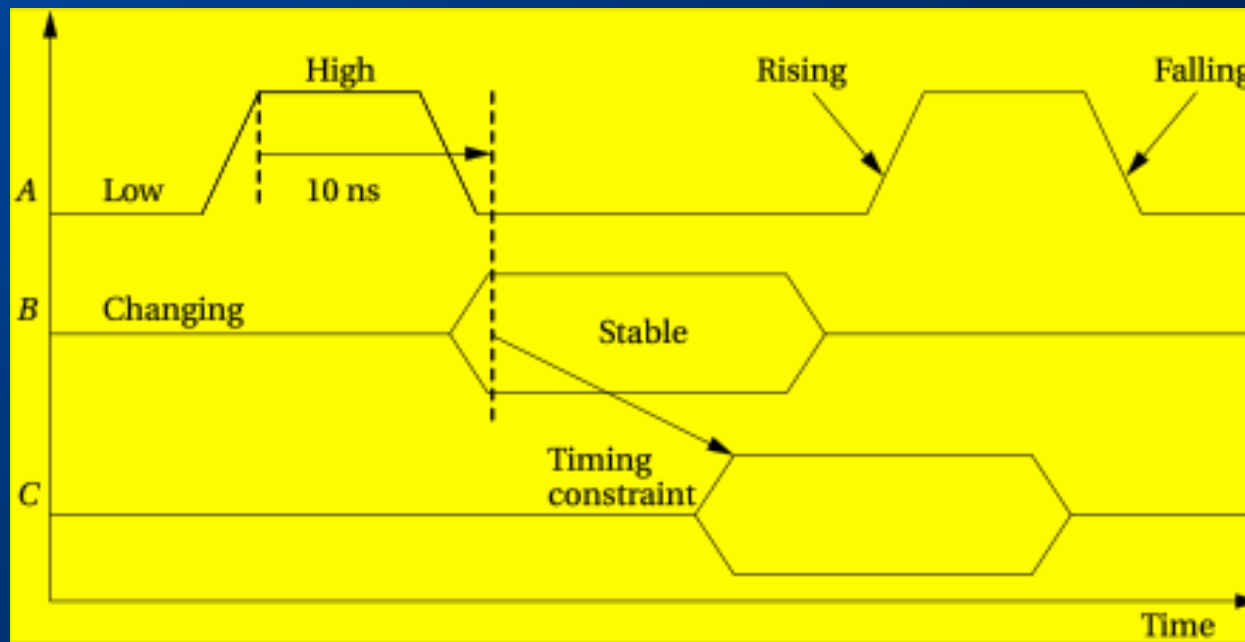


Microprocessor busses

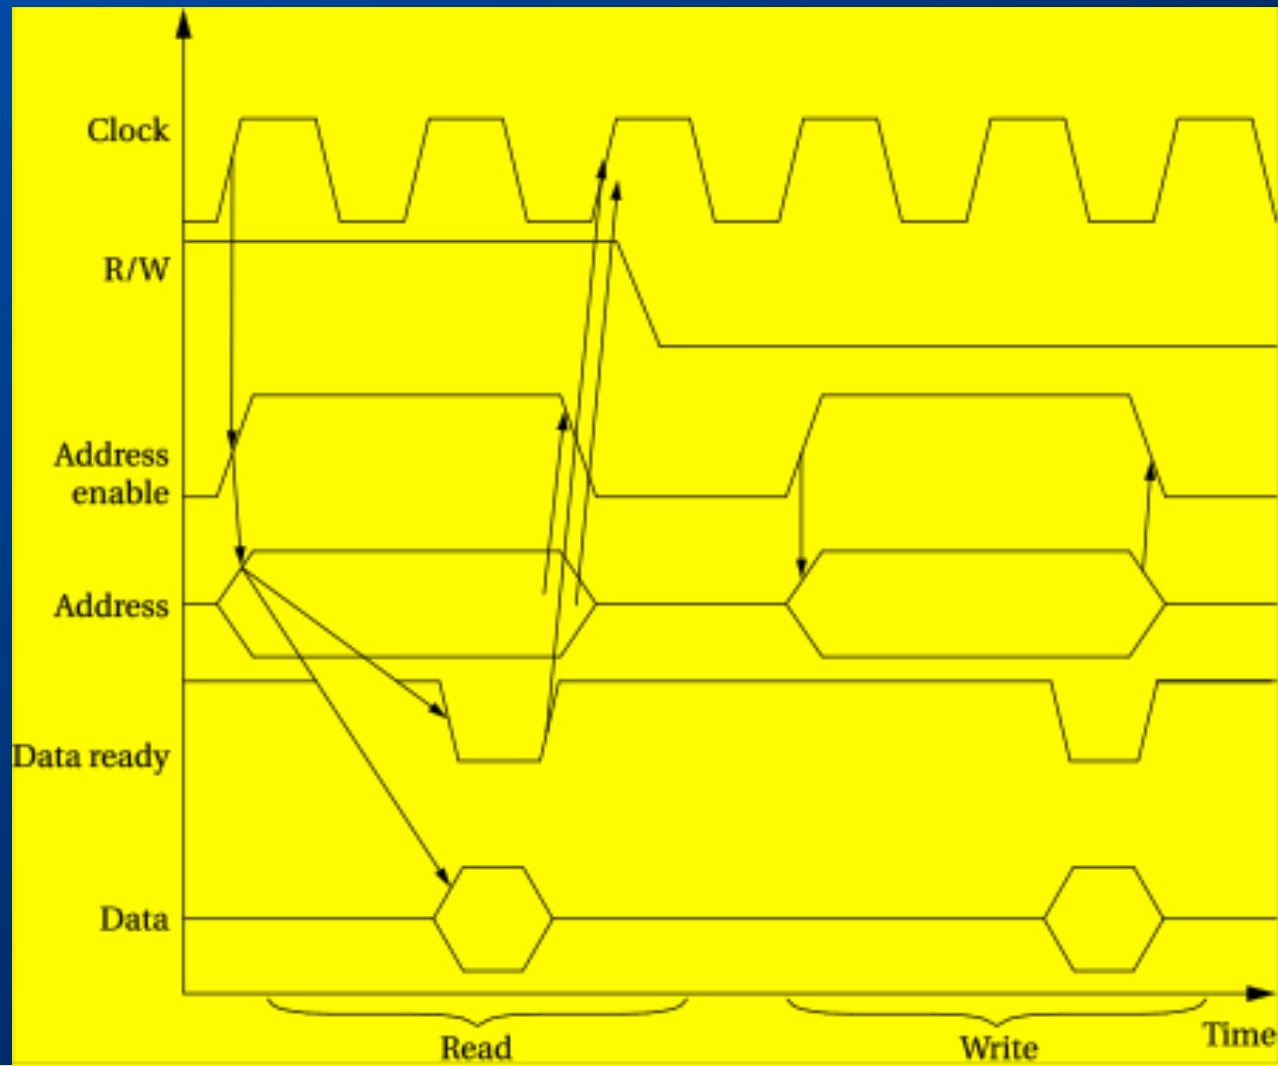
- Clock provides synchronization.
- R/W is true when reading (R/W' is false when reading).
- Address is a-bit bundle of address lines.
- Data is n-bit bundle of data lines.
- Data ready signals when n-bit data is ready.



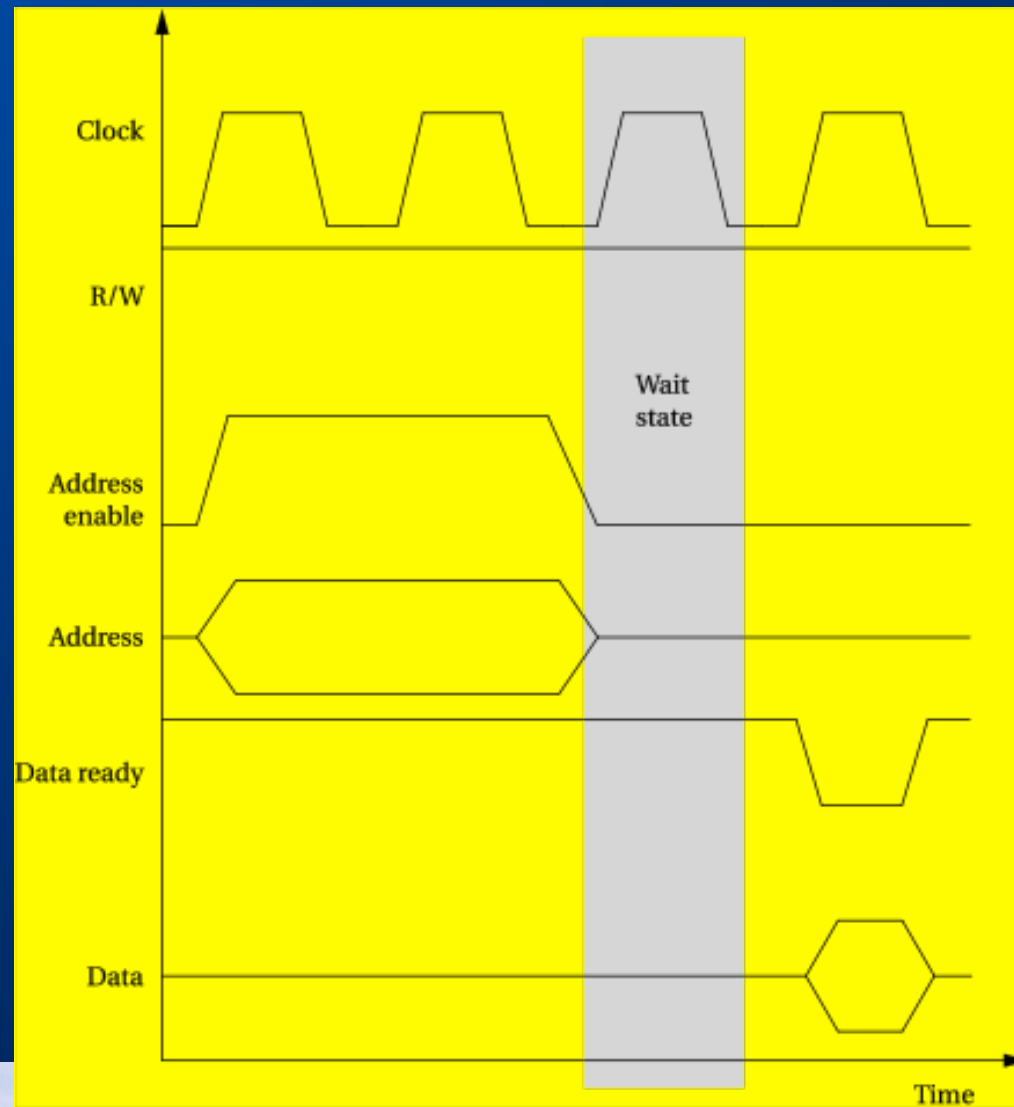
Timing diagrams



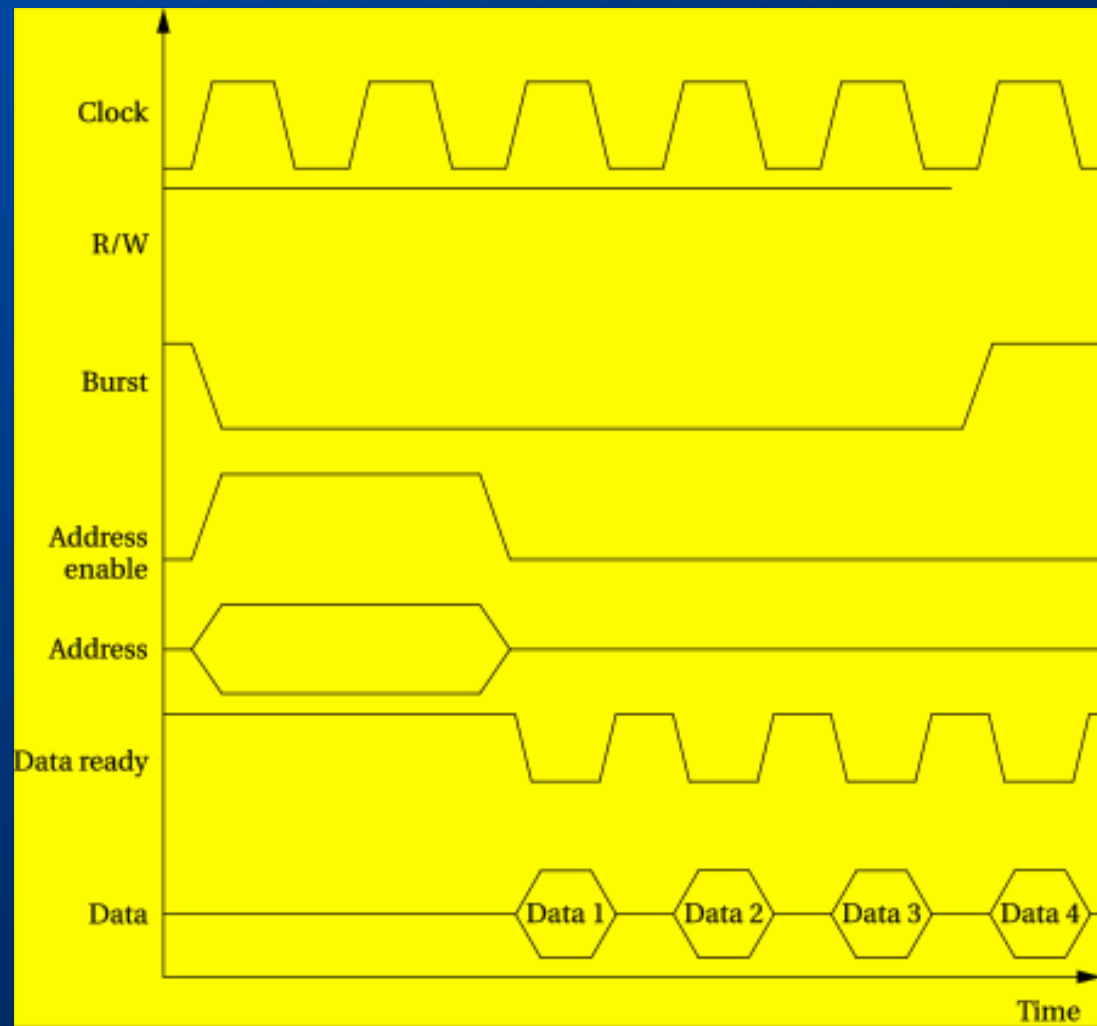
Bus read




Bus wait state



Bus burst read

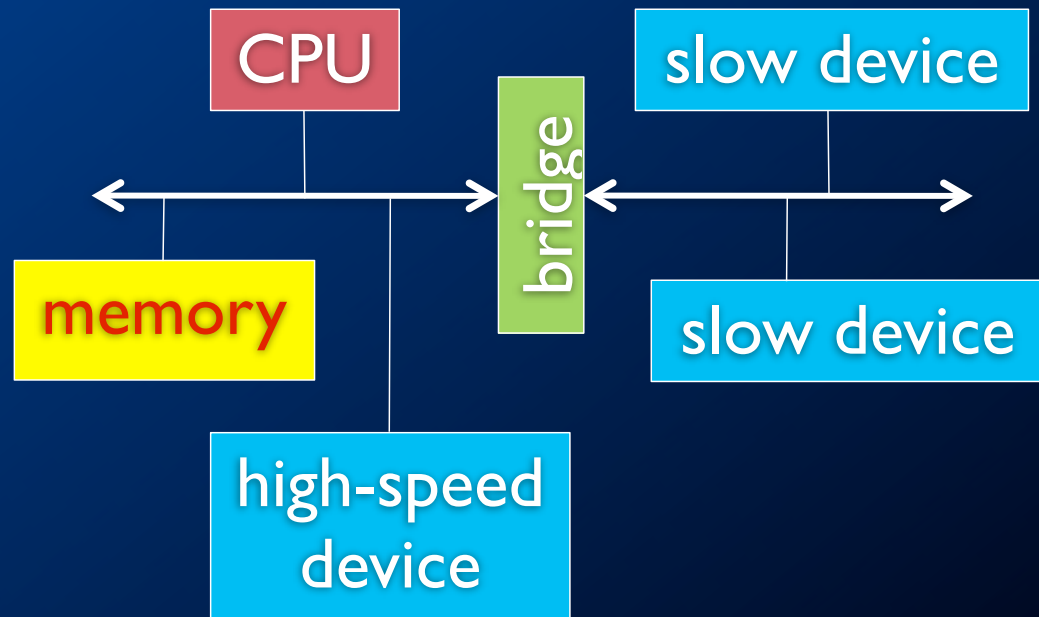


Bus multiplexing

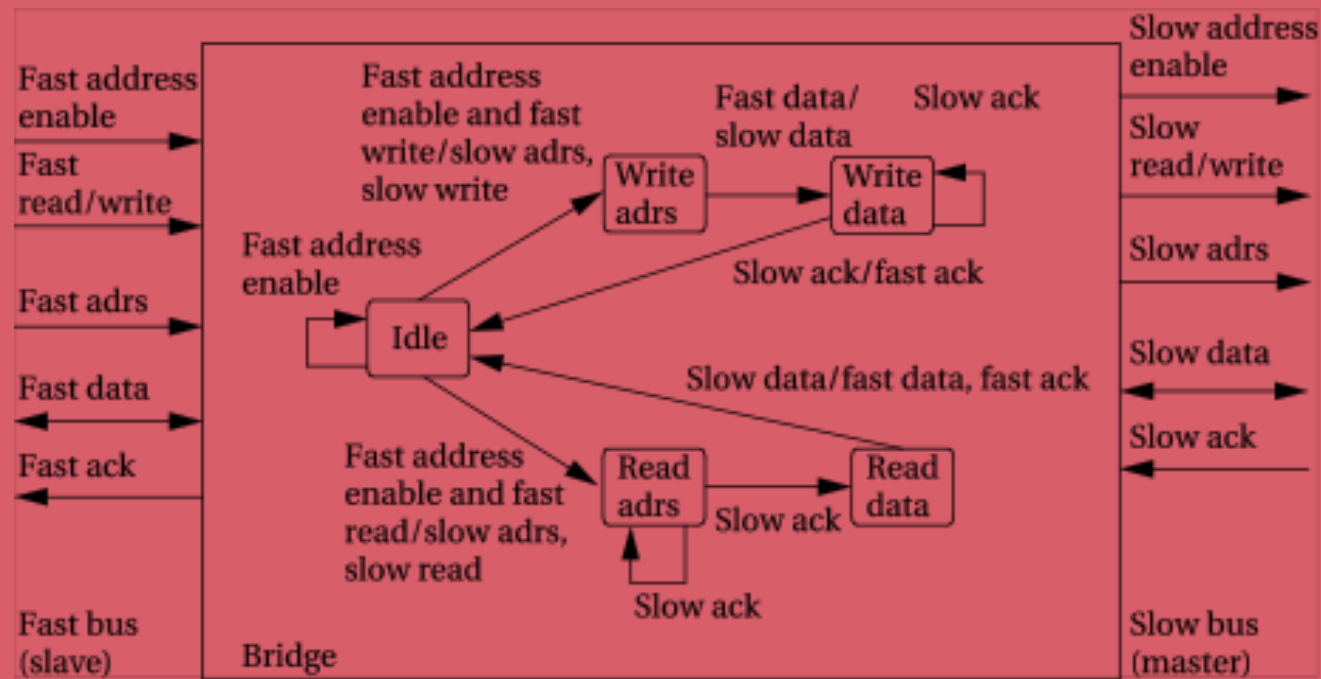


System bus configurations

- Multiple busses allow parallelism:
- Slow devices on one bus.
- Fast devices on separate bus.
- A bridge connects two busses.




Bridge state diagram



ARM AMBA bus

- Two varieties:
 - AHB is high-performance.
 - APB is lower-speed, lower cost.
- AHB supports pipelining, burst transfers, split transactions, multiple bus masters.
- All devices are slaves on APB.




Agenda

- Input and output
- Busses
- Memory Architectures



Memory components

- Several different types of memory:
 - DRAM.
 - SRAM.
 - Flash.
- Each type of memory comes in varying:
 - Capacities.
 - Widths.



Random-access memory

- Dynamic RAM is dense, requires refresh.
- Synchronous DRAM is dominant type.
- SDRAM uses clock to improve performance, pipeline memory accesses.
- Static RAM is faster, less dense, consumes more power.



Read-only memory

- ROM may be programmed at factory.
- Flash is dominant form of field-programmable ROM.
 - Electrically erasable, must be block erased.
 - Random access, but write/erase is much slower than read.
 - NOR flash is more flexible.
 - NAND flash is more dense.



Flash memory

- Non-volatile memory.
 - Flash can be programmed in-circuit.
- Random access for read.
- To write:
 - Erase a block to 1.
 - Write bits to 0.



Flash writing

- Write is much slower than read.
 - 1.6 μ s write, 70 ns read.
- Blocks are large (approx. 1 Mb).
- Writing causes wear that eventually destroys the device.
 - Modern lifetime approx. 1 million writes.



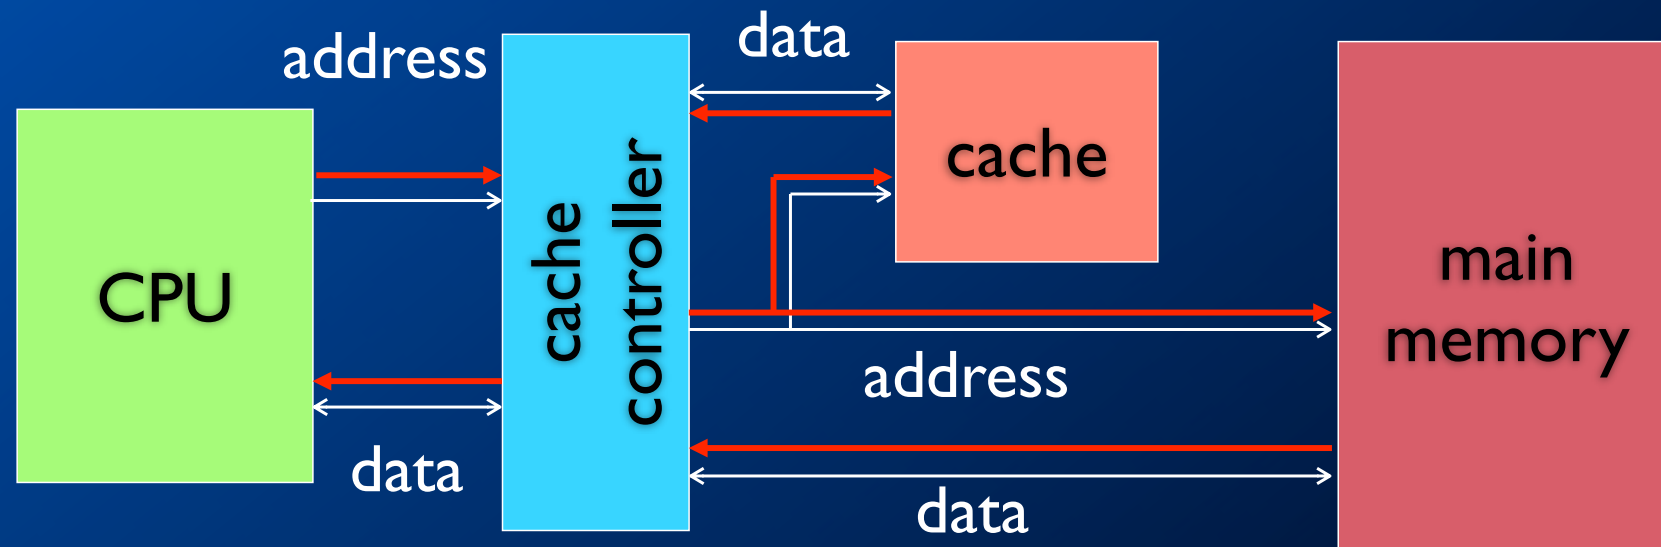
Types of flash

- NOR:
 - Word-accessible read.
 - Erase by blocks.
- NAND:
 - Read by pages (512-4K bytes).
 - Erase by blocks.
- NAND is cheaper, has faster erase, sequential access times.



	NOR	NAND
写入/擦除一个块的操作时间	1~5s	2~4ms
读性能	1200~1500KB	600~800KB
写性能	<80KB	200~400KB
接口/总线	SRAM接口/独立的地址数据总线	8位地址/数据/控制总线，I/O接口复杂
读取模式	随机读取	串行地存取数据
成本	较高	较低，单元尺寸约为NOR的一半，生产过程简单，同样大小的芯片可以做更大的容量
容量及应用场合	1~64MB，主要用于存储代码	8MB~4GB，主要用于存储数据
擦写次数(耐用性)	约10万次	约100万次
位交换(bit位反转)	少	较多，关键性数据需要错误探测/错误更正(EDC/ECC)算法
坏块处理	无，因为坏块故障率少	随机分布，无法修正

Caches and CPUs



Cache operation

- Many main memory locations are mapped onto one cache entry.
- May have caches for:
 - instructions;
 - data;
 - data + instructions (unified).
- Memory access time is no longer deterministic.



Terms

- Cache hit: required location is in cache.
- Cache miss: required location is not in cache.
- Working set: set of locations used by program in a time interval.



Types of misses

- Compulsory (cold): location has never been accessed.
- Capacity: working set is too large.
- Conflict: multiple locations in working set map to same cache entry.




Memory system performance

- h = cache hit rate.
- t_{cache} = cache access time, t_{main} = main memory access time.
- Average memory access time:
 - $t_{\text{av}} = ht_{\text{cache}} + (1-h)t_{\text{main}}$



Multiple levels of cache



Multi-level cache access time

- h_1 = cache hit rate.
- h_2 = hit rate on L2.
- Average memory access time:
 - $t_{av} = h_1 t_{L1} + h_2 t_{L2} + (1 - h_2 - h_1) t_{main}$



Replacement policies

- Replacement policy: strategy for choosing which cache entry to throw out to make room for a new memory location.
- Two popular strategies:
 - Random.
 - Least-recently used (LRU).



Cache organizations

- Fully-associative: any memory location can be stored anywhere in the cache (almost never implemented).
- Direct-mapped: each memory location maps onto exactly one cache entry.
- N-way set-associative: each memory location can go into one of n sets.

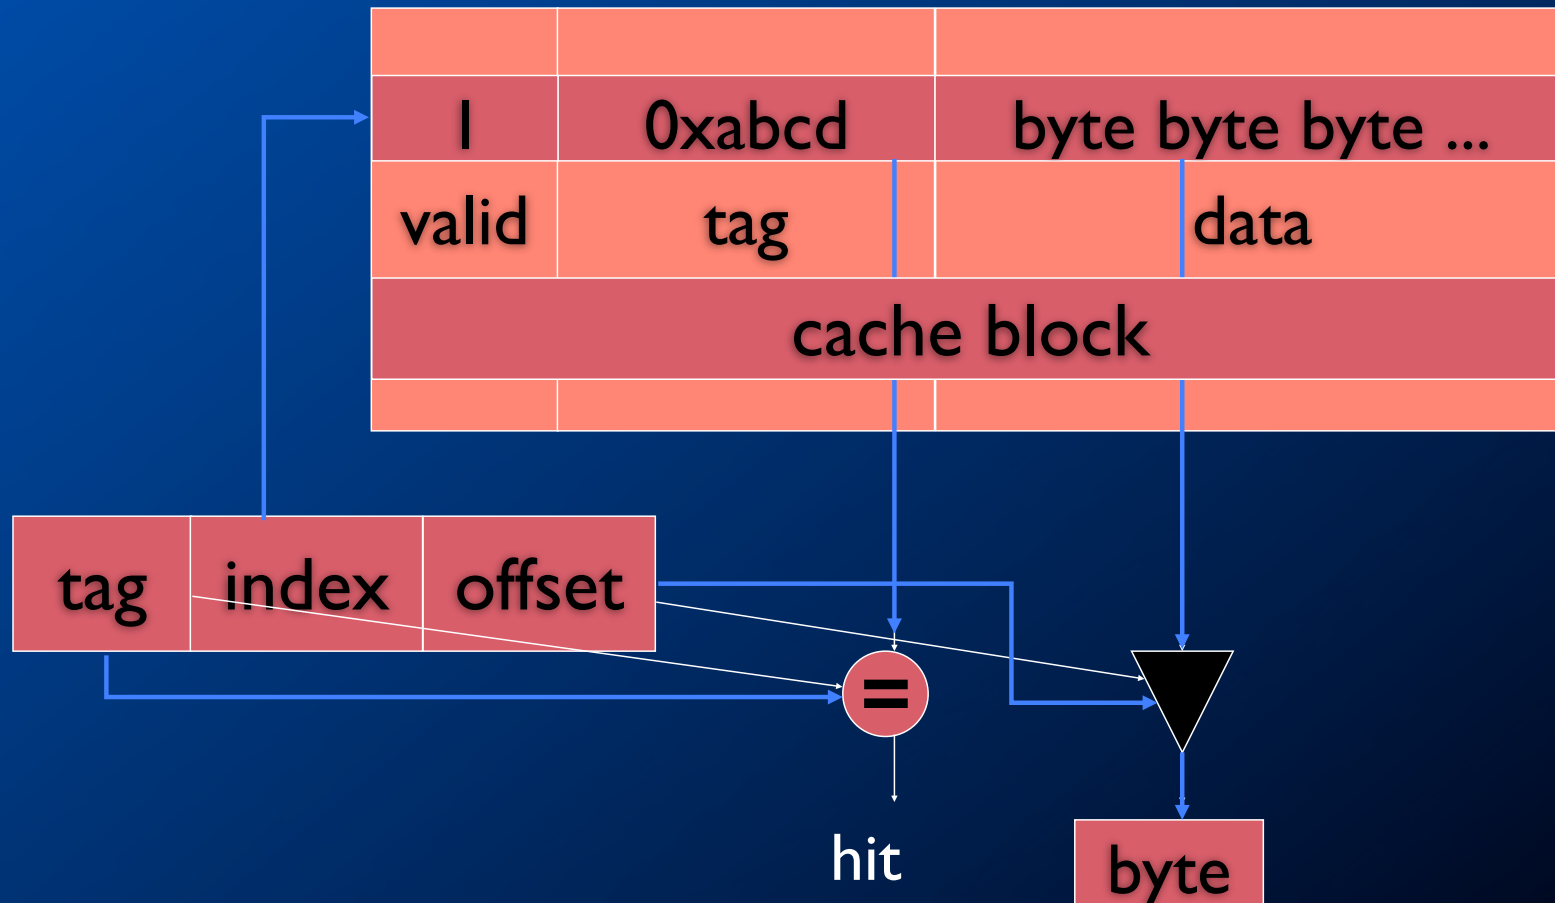


Cache performance benefits

- Keep frequently-accessed locations in fast cache.
- Cache retrieves more than one word at a time.
 - Sequential accesses are faster after first access.



Direct-mapped cache



Write operations

- Write-through: immediately copy write to main memory.
- Write-back: write to main memory only when location is removed from cache.




Direct-mapped cache locations

- Many locations map onto the same cache block.
- Conflict misses are easy to generate:
 - Array $a[]$ uses locations 0, 1, 2, ...
 - Array $b[]$ uses locations 1024, 1025, 1026, ...
 - Operation $a[i] + b[i]$ generates conflict misses.



Set-associative cache

- A set of direct-mapped caches:



Example: direct-mapped vs. set-associative

address	data
000	0101
001	1111
010	0000
011	0110
100	1000
101	0001
110	1010
111	0100



Direct-mapped cache behavior

- After 001 access:

block	tag	data
00	-	-
01	0	1111
10	-	-
11	-	-

- ♦ After 010 access:

block	tag	data
00	-	-
01	0	1111
10	0	0000
11	-	-



Direct-mapped cache behavior, cont'd.

◆ After 011 access:

block	tag	data
00	-	-
01	0	1111
10	0	0000
11	0	0110

◆ After 100 access:

block	tag	data
00	1	1000
01	0	1111
10	0	0000
11	0	0110



Direct-mapped cache behavior, cont'd.

◆ After 101 access:

block	tag	data
00	1	1000
01	1	0001
10	0	0000
11	0	0110

◆ After 111 access:

block	tag	data
00	1	1000
01	1	0001
10	0	0000
11	1	0100



2-way set-associative cache behavior

- ◆ Final state of cache (twice as big as direct-mapped):

Block	Way 0 tag	Way 0 data	Way 1 tag	Way 1 data
00	1	1000	-	-
01	0	1111	1	0001
10	0	0000	-	-
11	0	0110	1	0100



2-way set-associative cache behavior

- ◆ Final state of cache (same size as direct-mapped):

Block	Way 0 tag	Way 0 data	Way 1 tag	Way 1 data
0	01	0000	10	1000
1	10	0001	11	0100



Memory management units

- Memory management unit (MMU) translates addresses:



Memory management tasks

- Allows programs to move in physical memory during execution.
- Allows virtual memory:
 - memory images kept in secondary storage;
 - images returned to main memory on demand during execution.
- Page fault: request for location not resident in memory.



Address translation

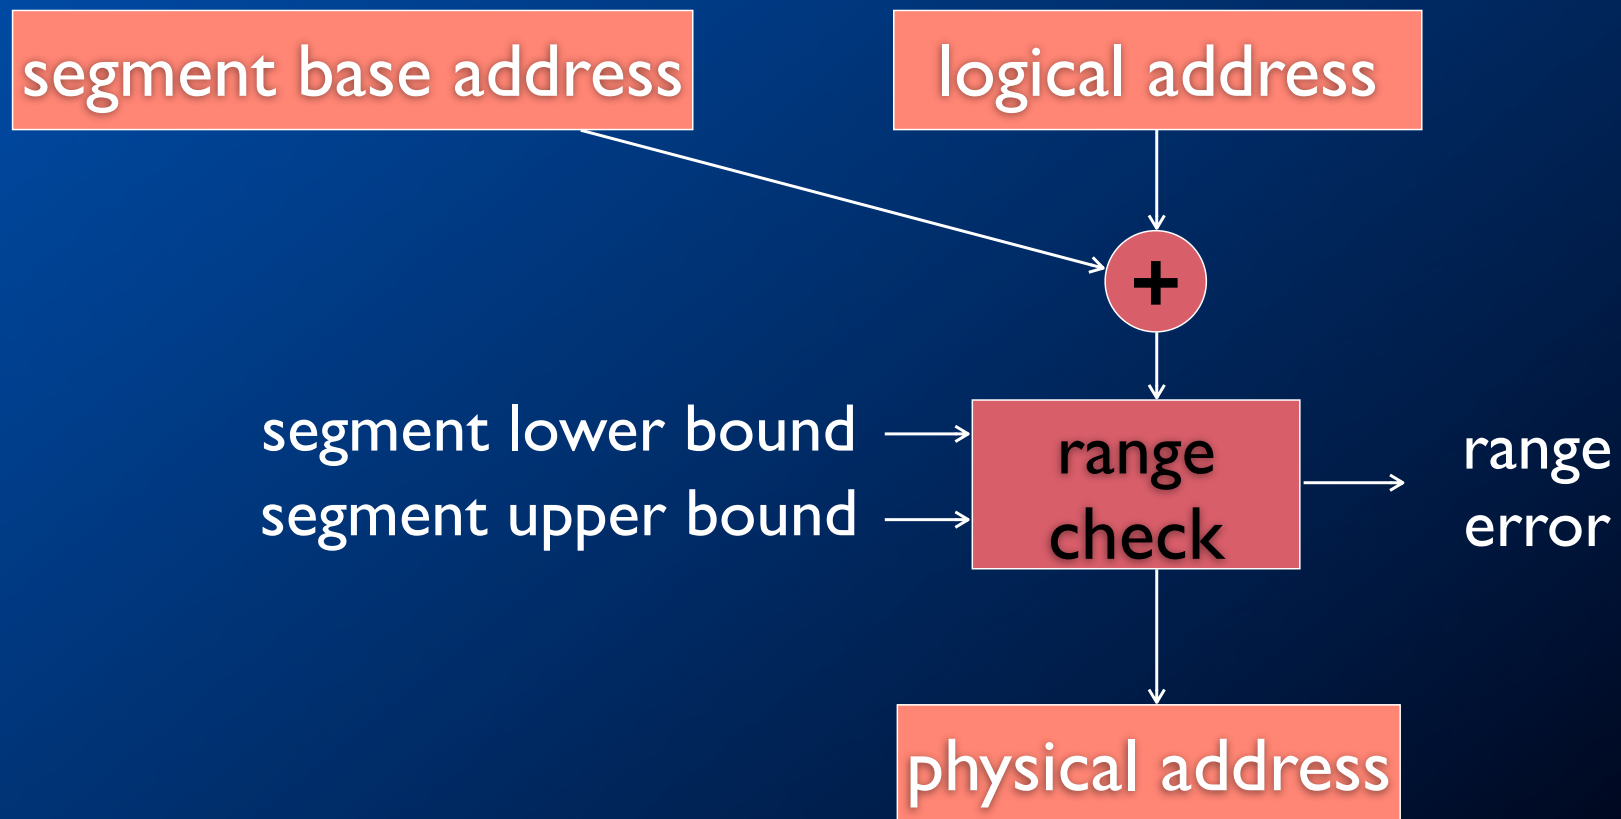
- Requires some sort of register/table to allow arbitrary mappings of logical to physical addresses.
- Two basic schemes:
 - segmented;
 - paged.
- Segmentation and paging can be combined (x86).



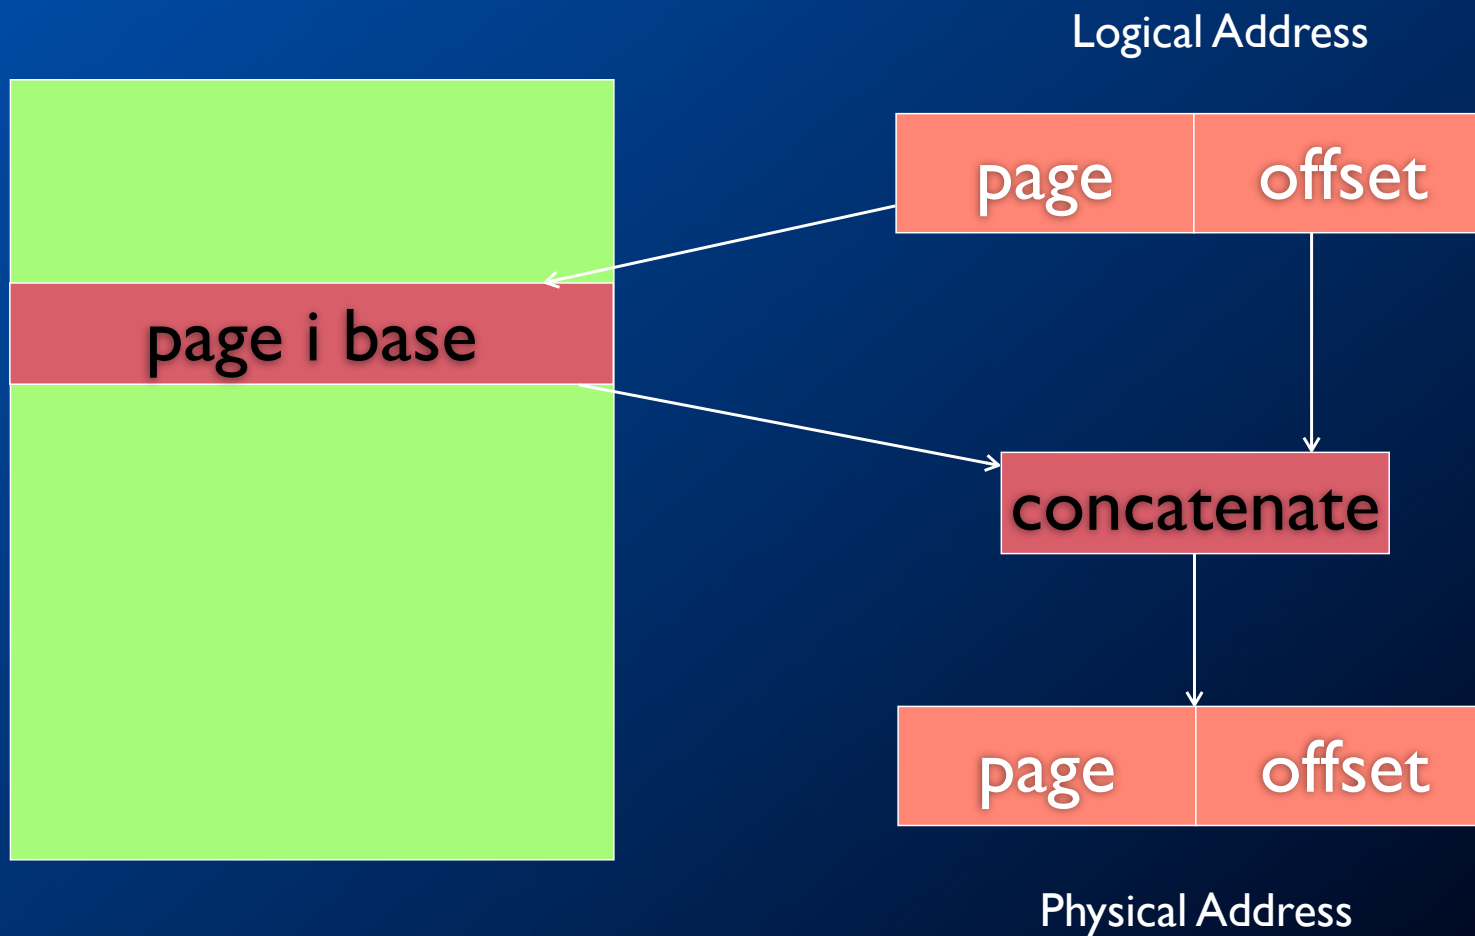
Segments and pages




Segment address translation




Page address translation



Page table organizations



flat



tree



Caching address translations

- Large translation tables require main memory access.
- TLB: cache for address translation.
 - Typically small.

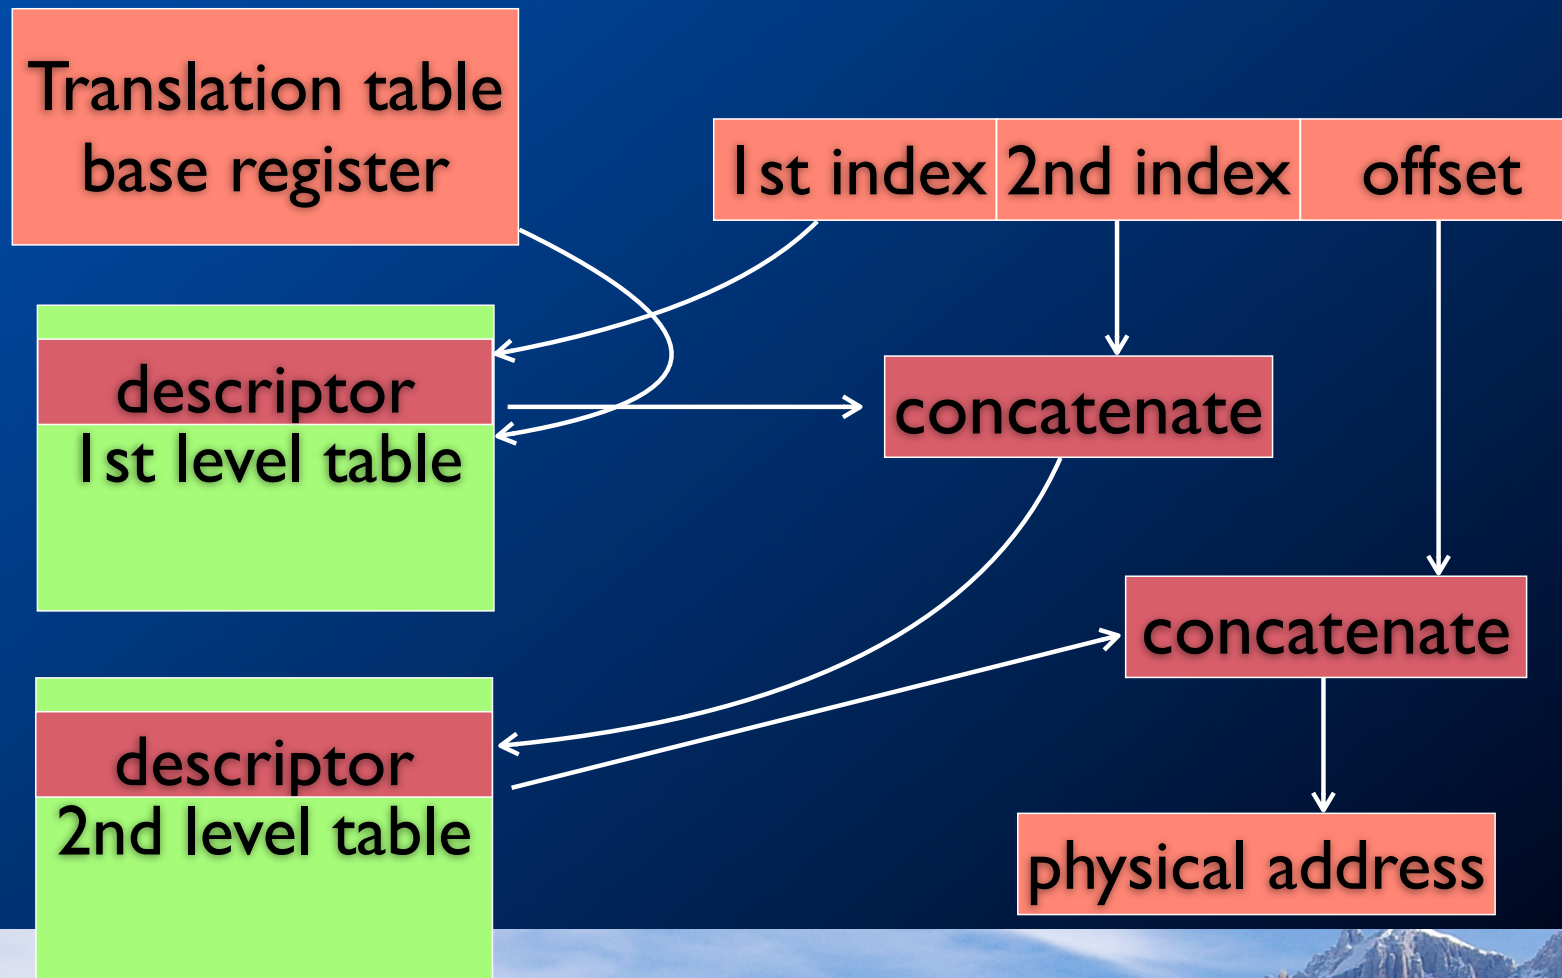


ARM memory management

- Memory region types:
 - section: 1 Mbyte block;
 - large page: 64 kbytes;
 - small page: 4 kbytes.
- An address is marked as section-mapped or page-mapped.
- Two-level translation scheme.



ARM address translation



Thank you!

